

All-Pairs Shortest Path using Fox's Algorithm

Project Overview

This project implements a parallel solution to the **All-Pairs Shortest Path Problem** using **Fox's Algorithm** for distributed matrix multiplication with **MPI (Message Passing Interface)**. The implementation uses the **Repeated Squaring Algorithm** with **min-plus matrix multiplication** to efficiently compute shortest paths between all pairs of nodes in a directed graph.

Problem Description

What the Project Does

The program solves the all-pairs shortest path problem for directed graphs by:

1. **Input:** Reading an adjacency matrix representing a directed graph where each element (i,j) represents the weight of the edge from node i to node j
2. **Processing:** Using Fox's algorithm to perform min-plus matrix multiplications in parallel across multiple MPI processes
3. **Output:** Producing a distance matrix where each element (i,j) represents the shortest path distance from node i to node j

The Problem It Solves

Given a directed graph $G = (V,E)$ with: - **V**: Set of vertices (nodes) - **E**: Set of edges (links) with weights

The goal is to find the shortest path distance between every pair of vertices. This is fundamental in: - **Network routing protocols** - **Transportation systems** - **Social network analysis** - **Game AI pathfinding** - **Supply chain optimization**

Parallel Computing Implementation

Fox's Algorithm

Fox's Algorithm is a parallel matrix multiplication algorithm designed for distributed memory systems. Key characteristics:

- **Process Grid:** Arranges P processes in a $\sqrt{P} \times \sqrt{P}$ grid
- **Matrix Partitioning:** Divides $N \times N$ matrices into $(N/\sqrt{P}) \times (N/\sqrt{P})$ blocks
- **Communication Pattern:** Each process communicates only with processes in the same row and column
- **Scalability:** Reduces communication overhead compared to naive parallel approaches

Min-Plus Matrix Multiplication

Instead of traditional matrix multiplication (\times , $+$), we use **min-plus algebra**: - **Multiplication operation** \rightarrow **Addition** - **Addition operation** \rightarrow **Minimum**

For matrices A and B, element (i,j) of result C:

$$C[i][j] = \min\{A[i][k] + B[k][j]\} \text{ for all } k$$

Repeated Squaring Algorithm

To find all shortest paths, we compute: - D_1 = initial adjacency matrix - $D_2 = D_1 \otimes D_1$ (shortest paths with ≤ 2 edges) - $D_4 = D_2 \otimes D_2$ (shortest paths with ≤ 4 edges) - ... - D_n = final result (shortest paths with $\leq N$ edges)

Where \otimes represents min-plus matrix multiplication.

Code Structure and Organization

Main Components

1. Main Function (main)

- **MPI Initialization:** Sets up MPI environment
- **Input Validation:** Checks Fox's algorithm constraints
- **Process Grid Setup:** Creates 2D Cartesian topology
- **Orchestration:** Coordinates the overall algorithm execution

2. Matrix Operations

- **allocate_matrix:** Dynamic memory allocation for 2D matrices
- **free_matrix:** Memory cleanup
- **read_input_matrix:** Input parsing with zero-to-infinity conversion
- **print_matrix:** Output formatting

3. Core Algorithm Functions

- **min_plus_multiply:** Performs min-plus multiplication on matrix blocks
- **min_plus_square:** Implements Fox's algorithm for min-plus matrix squaring
- **fox_algorithm:** Main Fox's algorithm implementation

4. Communication Setup

- **Grid Topology:** 2D Cartesian process grid with periodic boundaries
- **Row Communicators:** For broadcasting blocks within rows
- **Column Communicators:** For shifting blocks within columns

Data Structures

```
// Process grid coordinates
int grid_coords[2];           // [row, column] in process grid

// Communication contexts
MPI_Comm grid_comm;          // 2D Cartesian grid
MPI_Comm row_comm;           // Row-wise communication
MPI_Comm col_comm;           // Column-wise communication

// Matrix blocks (local to each process)
int **local_A;                // Input matrix block
int **local_result;           // Result matrix block
```

Algorithm Flow

1. Initialization

- Read matrix dimension and validate constraints
- Create 2D process grid topology
- Distribute matrix blocks to processes

2. Repeated Squaring Loop

```
for (int iter = 0; iter < ceil(log2(n)); iter++) {
    min_plus_square(local_result, local_result, ...);
}
```

3. Fox's Algorithm Steps (within each squaring operation)

- **Stage Loop:** For each stage (0 to $\sqrt{P} - 1$)
- **Broadcast:** Send appropriate A-block within row
- **Multiply:** Perform local min-plus multiplication
- **Shift:** Circularly shift B-blocks within column

4. Result Collection

- Gather all matrix blocks to process 0
- Reconstruct complete result matrix

Compilation and Execution Instructions

Prerequisites

- **MPI Implementation:** OpenMPI, MPICH, or Intel MPI
- **C Compiler:** GCC or Intel C Compiler
- **System:** Linux/Unix environment

Installing MPI (if not available) On Ubuntu/Debian systems:

```
sudo apt update
sudo apt install openmpi-bin openmpi-common openmpi-doc libopenmpi-dev
```

On CentOS/RHEL systems:

```
sudo yum install openmpi openmpi-devel
# or
sudo dnf install openmpi openmpi-devel
```

Note: If MPI is not available, the Makefile will automatically build a sequential version (`fox_sequential`) that validates the algorithm correctness using the Floyd-Warshall algorithm.

Compilation

```
# Using the provided Makefile
make
```

```
# Or manually
mpicc -Wall -Wextra -O3 -std=c99 fox.c -o fox -lm
```

Execution

Basic Usage

```
# Run with 4 processes
mpirun -np 4 ./fox < input.txt
```

```
# Run with 9 processes
mpirun -np 9 ./fox < input.txt
```

Input Format

```
6
0 2 0 5 0 0
0 0 0 0 0 0
0 2 0 0 0 5
0 0 0 0 1 0
3 9 3 0 0 0
0 0 0 0 1 0
```

Using Makefile Targets

```
# Create test input file
make test_input
```

```
# Run with different process counts
make run1    # 1 process
make run4    # 4 processes
```

```
make run9      # 9 processes
```

```
# Performance benchmark
```

```
make benchmark
```

Expected Outputs and Performance Analysis

Example Output

For the sample 6×6 graph input:

```
0 2 9 5 6 14
0 0 0 0 0 0
9 2 0 14 6 5
4 6 4 0 1 9
3 5 3 8 0 8
4 6 4 9 1 0
```

Performance Characteristics

Time Complexity

- **Sequential:** $O(N^3)$ for Floyd-Warshall
- **Parallel:** $O(N^3/P + \log N \times \text{communication_cost})$
- **Communication:** $O(\sqrt{P} \times N^2/P)$ per iteration

Space Complexity

- **Per Process:** $O(N^2/P)$ for local matrix blocks
- **Total:** $O(N^2)$ distributed across all processes

Expected Speedup For ideal conditions with P processes: - **Theoretical Maximum:** P (perfect scaling) - **Realistic:** $0.7P$ to $0.9P$ (considering communication overhead)

Performance Factors

1. **Matrix Size:** Larger matrices show better parallel efficiency
2. **Process Count:** Must be perfect squares (1, 4, 9, 16, 25, ...)
3. **Network Latency:** Affects communication-intensive phases
4. **Load Balance:** Fox's algorithm provides excellent load distribution

Benchmark Results Format

Execution time: X.XX ms

Process Count: P

Speedup: S.SS

Efficiency: E.EE%

Algorithms and Techniques Used

1. Fox's Algorithm

- **Purpose:** Parallel matrix multiplication for distributed memory
- **Key Idea:** 2D block decomposition with systematic communication
- **Advantage:** $O(\sqrt{P})$ communication complexity vs $O(P)$ for naive approaches

2. Min-Plus Algebra

- **Operations:** $(\min, +)$ semiring instead of $(+, \times)$ ring
- **Properties:** Associative, commutative (\min), distributive
- **Application:** Shortest path problems map naturally to min-plus algebra

3. Repeated Squaring

- **Concept:** Compute A, A^2, A^4, A^8, \dots until $A^{2^k} \geq A^N$
- **Efficiency:** $O(\log N)$ matrix multiplications vs $O(N)$ for naive approach
- **Convergence:** Result stabilizes when all shortest paths are found

4. MPI Cartesian Topology

- **Grid Creation:** `MPI_Cart_create` for 2D process arrangement
- **Subcommunicators:** Row and column communicators for efficient broadcasting
- **Coordinates:** Logical process positioning for algorithm coordination

5. Communication Patterns

- **Broadcast:** Row-wise distribution of A-blocks
- **Circular Shift:** Column-wise rotation of B-blocks
- **Point-to-Point:** Final result gathering

Error Handling and Constraints

Validation Checks

1. **Perfect Square Processes:** $P = Q^2$ where Q is integer
2. **Matrix Divisibility:** $N \bmod Q = 0$
3. **Input Format:** Proper matrix dimension and values
4. **Memory Allocation:** Robust error checking for malloc failures

Error Messages

Error: Number of processes (X) must be a perfect square

Error: Matrix dimension (N) must be divisible by sqrt(processes) (Q)

Error reading matrix element [i][j]

Memory allocation failed

Implementation Notes

Optimizations Applied

1. **Memory Layout:** Contiguous allocation for better cache performance
2. **Communication Overlap:** Asynchronous operations where possible
3. **Compiler Optimizations:** -O3 flag for aggressive optimization
4. **Infinity Handling:** Efficient representation using INT_MAX

Debugging Features

- **Timing Information:** Execution time measurement (excluding I/O)
- **Error Reporting:** Detailed error messages with context
- **Rank-based Output:** Process 0 handles all I/O operations

Scalability Considerations

- **Memory Per Process:** $O(N^2/P)$ scaling
- **Communication Volume:** $O(N^2/\sqrt{P})$ per process
- **Synchronization Points:** Minimized barrier operations

This implementation provides a robust, scalable solution to the all-pairs shortest path problem, demonstrating key concepts in parallel computing, distributed algorithms, and high-performance computing with MPI.